

QDMA Linux Kernel Reference Driver

User Guide

v2018.3 December 13, 2018



Notice of Disclaimer

The information disclosed to you here under (the "Materials") is provided solely for the selection and use of Xilinx products. To the maximum extent permitted by applicable law: (1) Materials are made available "AS IS" and with all faults, Xilinx hereby DISCLAIMS ALL WARRANTIES AND CONDITIONS, EXPRESS, IMPLIED, OR STATUTORY, INCLUDING BUT NOT LIMITED TO WARRANTIES OF MERCHANTABILITY, NON-INFRINGEMENT, OR FITNESS FOR ANY PARTICULAR PURPOSE; and (2) Xilinx shall not be liable (whether in contract or tort, including negligence, or under any other theory of liability) for any loss or damage of any kind or nature related to, arising under, or in connection with, the Materials (including your use of the Materials), including for any direct, indirect, special, incidental, or consequential loss or damage (including loss of data, profits, goodwill, or any type of loss or damage suffered as a result of any action brought by a third party) even if such damage or loss was reasonably foreseeable or Xilinx had been advised of the possibility of the same. Xilinx assumes no obligation to correct any errors contained in the Materials or to notify you of updates to the Materials or to product specifications. You may not reproduce, modify, distribute, or publicly display the Materials without prior written consent. Certain products are subject to the terms and conditions of the Limited Warranties which can be viewed at <http://www.xilinx.com/warranty.htm>; IP cores may be subject to warranty and support terms contained in a license issued to you by Xilinx.

© Copyright 2017-2018 Xilinx, Inc.

1 Revision History

Date	Version	Notes
18-Jun-2018	2018.1	Initial Version for Preliminary 2018.1 release
06-Aug-2018	2018.2	Updated the user guide for 2018.2 release: <ul style="list-style-type: none">• Added Appendix 1 to describe the dmactl command options in details
26-Sep-2018	2018.2.1	Updated for 2018.2.1 release <ul style="list-style-type: none">• Added sec 2.9, updated 2.8.3 for VF functionalities
13-Dec-2018	2018.3	Updated for 2018.3 release <ul style="list-style-type: none">• Updated module parameter options• Added Appendix 4 with doxygen tool usage

Table of Contents

1	Revision History	3
2	Introduction.....	6
2.1	<i>Document Overview.....</i>	<i>6</i>
2.2	<i>Document References</i>	<i>6</i>
2.3	<i>Glossary</i>	<i>6</i>
3	PCIe QDMA Driver for Linux Operating Systems.....	7
3.1	<i>Dependencies</i>	<i>7</i>
3.2	<i>Environment</i>	<i>8</i>
3.3	<i>Modifying the driver for your own PCIe device ID.....</i>	<i>8</i>
3.4	<i>Building the QDMA Driver Software.....</i>	<i>8</i>
3.5	<i>Installing the Compiled QDMA Driver binaries.....</i>	<i>9</i>
3.6	<i>Loading the QDMA Driver modules</i>	<i>9</i>
3.7	<i>Controlling and Configuring the QDMA IP</i>	<i>12</i>
3.8	<i>Running the VF on Virtual Machines</i>	<i>20</i>
3.9	<i>Un-installing the QDMA Driver modules</i>	<i>24</i>
4	Appendix 1 – User Application “dmactl” command options	25
5	Appendix 2 – dmautils tool	27
6	Appendix 3 – Release Directory Structure	31
6.1	<i>SW Directory</i>	<i>31</i>
7	Appendix 4 – Doxygen tool usage for document generation	32
7.1	<i>Steps for document generation</i>	<i>32</i>

LIST OF TABLES

<i>Table 1-1: Document References.....</i>	<i>6</i>
<i>Table 1-2: Glossary</i>	<i>6</i>
<i>Table 2-1: System Configuration</i>	<i>7</i>
<i>Table 2-2: QDMA Driver Supported Linux OS list</i>	<i>7</i>
<i>Table 4-1: dmautils tool configuration options</i>	<i>30</i>
<i>Table 5-1: SW Directory</i>	<i>31</i>

2 Introduction

2.1 Document Overview

The Xilinx PCI Express Multi Queue DMA (QDMA) IP provides high-performance direct memory access (DMA) via PCI Express. The PCIe QDMA can be implemented in UltraScale devices. This User Guide provides drivers and software that can be run on a PCI Express root port host PC to interact with the QDMA endpoint IP via PCI Express.

The drivers and software referenced in this User Guide are designed for Linux operating systems and can be used for lab testing or as a reference for driver and software development.

Through the use of the PCIe QDMA IP and the associated drivers and software you will be able to generate high-throughput PCIe memory transactions between a host PC and a Xilinx FPGA.

2.2 Document References

Document References	Version
[1] QDMA Subsystem for PCI Express (PG302)	2.0

Table 2-1: Document References

2.3 Glossary

Acronym / Term	Description
C2H	Card to Host
CLI	Command Line Interface
FPGA	Field Programmable Gate Array
H2C	Host to Card
IP	Intellectual Property
MM	Memory Mapped Mode
PF	Physical Function
QDMA	Multi Queue Direct Memory Access
ST	Streaming Mode
VF	Virtual Function

Table 2-2: Glossary

3 PCIe QDMA Driver for Linux Operating Systems

This User Guide document describes the following for QDMA Linux Driver that will be generally available for customers:

- Dependencies to be met for using the driver and environment to execute the driver
- Compiling and loading the driver
- Sample commands to use the driver

3.1 Dependencies

The release was tested with the following system configurations.

Directory		Description
Host System Configuration	Operating System	Ubuntu 16.04.3 LTS
	Linux Kernel	4.4.0-93-generic
	RAM	32GB
	Qemu version	QEMU emulator version 2.5.0 (Debian 1:2.5+dfsg-5ubuntu10.15)
Guest System (VM) Configuration	Operating System	Ubuntu 18.04 LTS
	Linux Kernel	4.15.1-20-generic
	RAM	4GB
	Cores	4

Table 3-1: System Configuration

Linux driver is supported on following OS and kernel versions.

Name	Version
CentOS	7.2.1511, 7.3.1611, 7.5.1804
RedHat	7.1
Ubuntu	16.04, 17.10.1, 18.04
Linux Kernel.org kernels	All long-term 3.16.56, 3.18.108, 4.1.51, 4.4.131, 4.9.99, 4.14.40, 4.15.18, 4.16.8

Table 3-2: QDMA Driver Supported Linux OS list

The following kernel functions shall be included in the OS kernel being used. Make sure that these functions are included in the kernel.

- Timer Functions
- PCIe Functions
- Kernel Memory functions

- Kernel threads
- Memory and GFP Functions

3.2 ***Environment***

To execute the QDMA driver on example design, following system requirements are to be met:

1. For best performance, use a Host System with at least one Gen 3 x16 PCIe slot and minimum 32GB RAM on same CPU node for 2K queues. For VM testing, host system must support virtualization and it must be enabled in the BIOS.
2. Any one of the Linux OS listed in Table 3-1
3. TULVU9P or VCU1525 FPGA Board
4. USB diligent cables to connect to the chosen board to the Host System.
5. Xilinx 2018.3 Vivado tools for programming the FPGA.

3.3 ***Modifying the driver for your own PCIe device ID***

During the PCIe DMA IP customization in Vivado you can specify a PCIe Device ID. This Device ID must be recognized by the driver in order to properly recognize the PCIe QDMA device. The current driver is designed to recognize the PCIe Device IDs that get generated with the PCIe example design when this value has not been modified. If you have modified the PCIe Device ID during IP customization you will need to modify the PCIe driver to recognize this new ID.

You may also want to modify the driver to remove PCIe Device IDs that will not be used by your solution. To modify the PCIe Device ID in the driver you should open the `drv/pci_ids.h` file and search for the `pcie_device_id` struct. This struct identifies the PCIe Device IDs that are recognized by the driver in the following format:

```
{PCI_DEVICE (0x10ee, 0x9034),},
```

Add, remove, or modify the PCIe Device IDs in this struct as desired for your application. The PCIe DMA driver will only recognize device IDs identified in this struct as PCIe QDMA devices. Once modified, the driver must be uninstalled and recompiled.

3.4 ***Building the QDMA Driver Software***

This driver supports both Physical Functions (PF) and Virtual Functions (VF).

In order to compile the Xilinx QDMA software, a configured and compiled Linux kernel source tree is required. The source tree may be only header files, or a complete tree. The source tree needs to be configured and the header files need to be compiled. And, the Linux kernel must be configured to use modules.

[Appendix 3](#) describes the Linux QDMA Driver software database structure and its contents on the Xilinx github (https://github.com/Xilinx/dma_ip_drivers, subdirectory QDMA/linux-kernel).

Compile the driver:

- `cd` into "QDMA/linux-kernel"

```
[xilinx@] # cd QDMA/linux-kernel
```


- Build the driver

```
[xilinx@] # make clean && make
```

From now on “linux-kernel” is assumed as the top-level directory and all the subsequent folders are mentioned relative this directory.

A sub-directory build/ will be created in “linux-kernel” after running "make". By default, both PF driver (qdma.ko) and VF driver (qdma_vf.ko) will be compiled along with the example application “dmactl”.

- If only PF driver is required, run make as

```
[xilinx@] # make pf
```

- If only VF driver is required, run make as

```
[xilinx@] # make vf
```

- If only example application needs to be compiled, run make as

```
[xilinx@] # make user
```

- For compiling the dma_to/from_device tools, run make as

```
[xilinx@] # make tools
```

3.5 ***Installing the Compiled QDMA Driver binaries***

To install the QDMA driver software, the installer must have the root permission.

Install the driver:

- Enter into “linux-kernel”

```
[xilinx@] # make install
```

- The QDMA module will be installed in /lib/modules/<linux_kernel_version>/updates/kernel/drivers/qdma directory.
- The “dmactl”, “dma_from_device” and “dma_to_device” tools will be installed in /user/local/sbin.

3.6 ***Loading the QDMA Driver modules***

Before loading the QDMA driver, make sure that an intended board is connected to the Host System and required bitstream is flashed on to the board.

Load the QDMA driver:

QDMA driver can be loaded in poll mod, direct interrupt mode or indirect interrupt mode. QDMA driver supports the following module parameters.

Module Parameter Name	Description
mode	<p>mode module parameter is used to enable the qdma driver functionality in different modes.</p> <p>Kernel module can be loaded in following different modes</p> <ul style="list-style-type: none"> 0 - Auto Mode, driver decides to process the request in poll or interrupt mode 1 - Poll Mode 2 - Direct Interrupt Mode 3 - Interrupt Aggregation Mode or Indirect Interrupt Mode 4 - Legacy Interrupt Mode <p>By default, mode is set to 3 and driver is loaded in indirect interrupt mode</p> <p>To load the driver in poll mode, use the below command.</p> <p>Ex: <code>insmod qdma.ko mode=1</code></p> <p>To load the driver in direct interrupt mode, use the below command.</p> <p>Ex: <code>insmod qdma.ko mode=2</code></p> <p>To load the driver in indirect interrupt mode, use the below command.</p> <p>Ex: <code>insmod qdma.ko mode=3</code></p>
master_pf	<p>master_pf module parameter is used to set the master pf for qdma driver</p> <p>By default, master_pf is set to PF0(First device in the PF list)</p> <p>To set any other PF as master_pf, use the module parameter as below</p> <p><code>insmod qdma.ko master_pf=<pf_bdf_number></code></p> <p><i>lspci grep Xilinx</i></p> <p>01:00.1 Memory controller: Xilinx Corporation Device 913f</p> <p>Ex: <code>insmod qdma.ko master_pf=0x01001</code></p>

	<p>When multiple cards are inserted in the same host system and master_pf needs to be updated for each card, use the command as below.</p> <p><i>lspci grep Xilinx</i></p> <p>01:00.1 Memory controller: Xilinx Corporation Device 913f</p> <p><i>lspci grep Xilinx</i></p> <p>02:00.1 Memory controller: Xilinx Corporation Device 913f</p> <p>Ex: insmod qdma.ko master_pf=0x01001, 0x02001</p>
tm_mode_en	<p>tm_mode_en parameter is used to enable Traffic Manager mode in driver to test desc bypass functionality with Traffic Manager example design for ST H2C queue.</p> <p>By default, tm_mode_en is set to 0.</p> <p>To load driver with Traffic Manager mode enabled, use below command:</p> <p>Ex. insmod qdma.ko tm_mode_en=1</p> <p>NOTE: <i>This parameter is experimental and should only be used only with Traffic Manager example design.</i></p>
tm_one_cdh_en	<p>tm_one_cdh_en is used to test 1 CDH (Custom Defined Header) functionality with Traffic Manager example design when driver is loaded with tm_mode_en set to 1.</p> <p>By default, tm_one_cdh_en is set to 0 indicating that driver will send pkts with Zero CDH.</p> <p>To load driver with 1 CDH enabled, use below command:</p> <p>Ex. insmod qdma.ko tm_mode_en=1 tm_one_cdh_en=1</p> <p>NOTE: <i>This parameter is experimental and should only be used only with Traffic Manager example design.</i></p>

- Load the driver in poll mode as

```
[xilinx@] # modprobe qdma mode=1
```

- Load the driver in direct interrupt mode as

```
[xilinx@] # modprobe mode=2
```

- Load the driver in indirect interrupt mode as

```
[xilinx@] # modprobe qdma mode=3
```

- Load the driver on a VM

Auto mode: [xilinx@] # modprobe qdma_vf mode=0

Poll mode: [xilinx@] # modprobe qdma_vf mode=1

Direct interrupt mode: [xilinx@] # modprobe qdma_vf mode=2

Indirect interrupt mode: [xilinx@] # modprobe qdma_vf mode=3

Now the QDMA software is ready for use.

3.7 *Controlling and Configuring the QDMA IP*

3.7.1 Configuration through sysfs

Once the qdma pf module is inserted and until any queue is added into the system and FMAP programming is not done, sysfs provides an interface to configure some parameters for the module configuration.

```
[xilinx@] # lspci | grep -i Xilinx
81:00.0 Memory controller: Xilinx Corporation Device 903f
81:00.1 Memory controller: Xilinx Corporation Device 913f
81:00.2 Memory controller: Xilinx Corporation Device 923f
81:00.3 Memory controller: Xilinx Corporation Device 933f
```

Based on the above lspci output, traverse to “/sys/bus/pci/devices/<device node>/qdma” to find the list of configurable parameters for each PF.

Below table describes the various configurable parameters through sysfs.

Parameter name	Description	Example
qmax	Maximum number of queues associated for the current pf are displayed here. Currently 2048 queues are dedicated for PFs and each	Display the current value: [xilinx@] # cat /sys/bus/pci/devices/0000:81:00.0/qdma/qmax

	<p>PF gets 512 queues each by default.</p> <p>If the queue allocation needs to be different for any PF, access the qmax sysfs entry and set the required number.</p> <p>Once the number of queues for any PF is changed from the default value, the remaining set of queues among the 2048 queues are evenly distributed for the remaining PFs.</p>	<p>Set a new value:</p> <pre>[xilinx@] # echo 1024 > /sys/bus/pci/devices/0000:81:00.0/qdma/qmax</pre> <p>Ex: Default queue sets for all PFs</p> <pre>[xilinx@] # dmactl dev list</pre> <pre>qdma81000 0000:01:00.0 max QP: 449, 0~448 qdma81001 0000:01:00.1 max QP: 449, 449~897 qdma81002 0000:01:00.2 max QP: 449, 898~1346 qdma81003 0000:01:00.3 max QP: 449, 1347~1795</pre> <pre>xilinx@] #echo 1770 > /sys/bus/pci/devices/0000:81:00.0/qdma/qmax</pre> <pre>[xilinx@] # dmactl dev list</pre> <pre>qdma81000 0000:01:00.0 max QP: 1770, 0~1769 qdma81001 0000:01:00.1 max QP: 8, 1770~1777 qdma81002 0000:01:00.2 max QP: 8, 1778~1785 qdma81003 0000:01:00.3 max QP: 8, 1786~1793</pre>
qmax_vf	<p>QDMA IP supports 2048 queues and all the queues are allocated to PFs by default.</p> <p>qmax_vf sysfs entry is used to allocate the queues to VF. This entry is available only for master_pf.</p> <p>Before instantiating the VFs, allocate required number of queues for VFs from the available pool.</p>	<p>Assume that PF0 is the master PF.</p> <p>Display the current value:</p> <pre>[xilinx@] # cat /sys/bus/pci/devices/0000:81:00.0/qdma/qmax_vfs</pre> <p>Set a new value:</p> <pre>[xilinx@] # echo 1024 > /sys/bus/pci/devices/0000:81:00.0/qdma/qmax_vfs</pre>
intr_rngsz	<p>Interrupt ring size is associated with indirect interrupt mode.</p> <p>When the module is inserted in indirect interrupt mode, by default the interrupt aggregation ring size is set 0 i.e 512 entries</p> <p>User can configure the interrupt ring entries in multiples of 512 hence set the intr_ring_size with multiplication factor</p>	<p>Display the current value:</p> <pre>[xilinx@] # cat /sys/bus/pci/devices/0000:81:00.0/qdma/intr_rngsz</pre> <p>Set a new value:</p> <pre>[xilinx@] # echo 2 > /sys/bus/pci/devices/0000:81:00.0/qdma/intr_rngsz</pre>

	<p>0 - INTR_RING_SZ_4KB, Accommodates 512 entries</p> <p>1 - INTR_RING_SZ_8KB, Accommodates 1024 entries</p> <p>2 - INTR_RING_SZ_12KB, Accommodates 1536 entries</p> <p>3 - INTR_RING_SZ_16KB, Accommodates 2048 entries</p> <p>4 - INTR_RING_SZ_20KB, Accommodates 2560 entries</p> <p>5 - INTR_RING_SZ_24KB, Accommodates 3072 entries</p> <p>6 - INTR_RING_SZ_24KB, Accommodates 3584 entries</p> <p>7 - INTR_RING_SZ_24KB, Accommodates 4096 entries</p>	
wrb_acc	<p>Completion interval if Completions are enabled for a queue configured for internal mode.</p> <p>3'h0: 4</p> <p>3'h1: 8</p> <p>3'h2: 16</p> <p>3'h3: 32</p> <p>3'h4: 64</p> <p>3'h5: 128</p> <p>3'h6: 256</p> <p>3'h7: 512</p> <p>Completion accumulation value is calculated as $2^{\text{(register bit [2:0])}}$. Maximum accumulation is 512.</p> <p>Accumulation can be disabled via queue context</p>	<p>Display the current value:</p> <pre>[xilinx@] # cat /sys/bus/pci/devices/0000:81:00.0/qdma/ wrb_acc</pre> <p>Set a new value:</p> <pre>[xilinx@] # echo 2 > /sys/bus/pci/devices/0000:81:00.0/qdma/ wrb_acc</pre>
sriov_numvfs	<p>QDMA IP supports 252 VFs.</p> <p>Identify the number of VFs supported for each PF using the sriov_totalvfs sysfs entry.</p>	<p>Assume that PF0 is the master PF.</p> <p>Display the currently supported max VFs:</p> <pre>[xilinx@] # cat /sys/bus/pci/devices/0000:81:00.0/sriov_totalvfs</pre> <p>Instantiate the required number of VFs for a PF:</p> <pre>[xilinx@] # echo 3 > /sys/bus/pci/devices/0000:81:00.0/sriov_numvfs</pre>

		<p>Once the VFS are instantiated, required number of queues can be allocated the VF using qmax sysfs entry available in VF at</p> <p><code>/sys/bus/pci/devices/<VF function number>/qdma/qmax</code></p>
--	--	---

3.7.2 Control and configuration through “dmactl”

QDMA driver comes with a command-line configuration utility called “dmactl” to manage the driver.

The Xilinx QDMA control tool, dmactl, is a Command Line utility which is installed in `/usr/local/sbin/` and allows administration of the Xilinx QDMA queues. Make sure that the installation path “`/usr/local/sbin/`” is added to the “PATH” variable.

It can perform the following functions:

- Query the QDMA functions/devices the driver has bind into
- Query control and configuration
 - ✓ List all the queues on a device/function
 - ✓ Add/configure a new queue on a device/function
 - ✓ Start an already added/configured queue (i.e., bring the queue online)
 - ✓ Stop a started queue (i.e., bring the queue offline)
 - ✓ Delete an already added/configured queue
- register access
 - ✓ Read a register
 - ✓ Write a register
 - ✓ Dump the qdma config bar and user bar registers
- debug helper
 - Display a queue's configuration parameters
 - Display a queue's descriptor ring entries
 - Display a c2h queue's completion ring entries
 - Display the interrupt ring entries
- For help run
 - `dmactl -h`

For more details on the dmactl tool commands and options for each command, refer to dmactl man page.

- For dmactl man page, run
 - `man dmactl`

dma_to_device: This utility is used to transfer the data from Host to Card(H2C). It requires input as the name of the device node and the size of the transfer as mandatory parameters. User “`dma_to_device -help`” to see the various options supported for this utility.

dma_from_device: This utility is used to transfer the data from Card to Host(C2H). It requires input as the name of the device node and the size of the transfer as mandatory parameters. User “dma_from_device -help” to see the various options supported for this utility.

3.7.2.1 Example: Get the list of devices the driver has bind with

List the devices using lspci to cross check the devices are detected as PCIe devices

```
[xilinx@] # lspci | grep -i Xilinx
81:00.0 Memory controller: Xilinx Corporation Device 903f
81:00.1 Memory controller: Xilinx Corporation Device 913f
81:00.2 Memory controller: Xilinx Corporation Device 923f
81:00.3 Memory controller: Xilinx Corporation Device 933f
```

```
[xilinx@] # dmactl dev list
```

```
qdma81000 0000:01:00.0 max QP: 448, 0~447
qdma81001 0000:01:00.1 max QP: 448, 512~959
qdma81002 0000:01:00.2 max QP: 448, 1024~1471
qdma81003 0000:01:00.3 max QP: 448, 1536~1983
```

3.7.2.2 Example: Configure and control a queue in Memory Mapped(MM) Mode

- ✓ Add a queue on qdma0

```
[root@] # dmactl qdma81000 q add idx 0 mode mm dir h2c
```

```
qdma81000: 01:00.00 config bar: 0, user bar: 2, max #. QP: 448
```

```
qdma81000 -MM-0 H2C added.
Added 1 Queues.
```

Note: Change the dir to “c2h” for Card-to-Host direction

- ✓ Start an already added queue

```
[root@] # dmactl qdma81000 q start idx 0 dir h2c
```

```
qdma81000: 01:00.00 config bar: 0, user bar: 2, max #. QP: 448
```

```
Started Queues 0 -> 0.
```

Note: Change the dir to “c2h” for Card-to-Host direction

*After the queue is started the normal read and write operation can be performed on the character device /dev/qdma81000-MM-C2H-0.

- ✓ Perform a dma transfer from Host to Card (H2C)


```
[root@] # dma_to_device -d /dev/qdma81000-MM-0 -s 512
** Average BW = 512, 4.289041
```

- ✓ Perform a dma transfer from or Card to Host (C2H)

```
[root@] # dma_from_device -d /dev/ qdma81000-MM-0 -s 512
** Average BW = 512, 4.289041
```

- ✓ Stop a queue

```
[root@] # dmactl qdma81000 q stop idx 0 dir h2c
qdma81000: 01:00.00      config bar: 0, user bar: 2, max #: QP: 448
```

Stopped Queues 0 -> 0.

Note: Change the dir to "c2h" for Card-to-Host direction

- ✓ Delete a queue

```
[root@] # dmactl qdma81000 q del idx 0 dir h2c
qdma81000: 01:00.00      config bar: 0, user bar: 2, max #: QP: 448
```

Deleted Queues 0 -> 0.

Note: Change the dir to "c2h" for Card-to-Host direction

3.7.2.3 Example: Configure and control a queue in Streaming (ST) H2C Mode

- ✓ Add a queue on qdma0

```
[root@] # dmactl qdma81000 q add idx 0 mode st dir h2c
qdma81000: 01:00.00      config bar: 0, user bar: 2, max #: QP: 448
```

qdma81000-ST-0 H2C added.
Added 1 Queues.

- ✓ Start an already added queue

```
[root@] # dma dmactl qdma81000 q start idx 0 dir h2c
qdma81000: 01:00.00      config bar: 0, user bar: 2, max #: QP: 448
```

Started Queues 0 -> 0.

- ✓ Perform a dma transfer from Host-to-Card (H2C)

```
[root@] # dma_to_device -d /dev/ qdma81000-ST-0 -s 512
** Average BW = 512, 4.289041
```

- ✓ Stop a queue

```
[root@] # dmactl qdma81000 q stop idx 0 dir h2c
qdma81000: 01:00.00      config bar: 0, user bar: 2, max #: QP: 448
```

Stopped Queues 0 -> 0.

Note: Change the dir to "c2h" for Card-to-Host direction

- ✓ Delete a queue

```
[root@] # dmactl qdma81000q del idx 0 dir h2c
qdma81000: 01:00.00      config bar: 0, user bar: 2, max #: QP: 448
```

Deleted Queues 0 -> 0.

3.7.2.4 Example: Configure and control a queue in Streaming (ST) C2H Mode

NOTE: the following example with user bar register access is based on the Streaming Mode (ST) example design.

- ✓ Add a MM H2C queue on qdma81000

```
[root@] # dmactl qdma81000 q add idx 0 mode st dir h2c
qdma81000: 01:00.00      config bar: 0, user bar: 2, max #: QP: 448
```

qdma81000-ST-0 C2H added.
Added 1 Queues.

- ✓ Start an already added queue

```
[root@] # dma dmactl qdma81000 q start idx 0 dir c2h
qdma81000: 01:00.00      config bar: 0, user bar: 2, max #: QP: 448
```

Started Queues 0 -> 0.

- ✓ Write the HW Qid number in user bar register 0x0

```
[root@] # dmactl qdma81000 reg write bar 2 0x0 0
qdma0: 01:00.00      config bar: 0, user bar: 2, max #: QP: 448
qdma0, 01:00.00, bar#2, reg 0x0 -> 0x0, read back 0x0.
```

- ✓ Program the size and number of packets in to user bar registers 0x4 and 0x20 respectively

```
[root@] # dmactl qdma81000 reg write bar 2 0x4 512
qdma81000: 01:00.00      config bar: 0, user bar: 2, max #: 448
qdma81000, 01:00.00, bar#2, reg 0x4 -> 0x200, read back 0x200.
```

```
[root@] # dmactl qdma81000 reg write bar 2 0x20 1
qdma81000: 01:00.00      config bar: 0, user bar: 2, max #: 448
qdma81000, 01:00.00, bar#2, reg 0x20 -> 0x1, read back 0x1.
```

- ✓ Perform a dma transfer from Card-to-Host (C2H)

```
[root@] # dma_from_device -d /dev/ qdma81000-ST-0 -s 512
** Average BW = 512, 4.289041
```

- ✓ Stop a queue

```
[root@] # dmactl qdma81000 q stop idx 0 dir h2c
qdma81000: 01:00.00      config bar: 0, user bar: 2, max #: 448
```

Stopped Queues 0 -> 0.

Note: Change the dir to “c2h” for Card-to-Host direction

- ✓ Delete a queue

```
[root@] # dmactl qdma81000 q del idx 0 dir h2c
qdma81000: 01:00.00      config bar: 0, user bar: 2, max #: 448
```

Deleted Queues 0 -> 0.

3.7.3 Adding VFs to PFs

This section provides the details on assigning VFs to the PFs.

Design supports 252 VFs in total and based on the HW design, VFs can be assigned to PFs as below using sysfs

- ✓ Display the current available VFs for a PF:

```
[xilinx@] # lspci | grep -i Xilinx
81:00.0 Memory controller: Xilinx Corporation Device 903f
81:00.1 Memory controller: Xilinx Corporation Device 913f
81:00.2 Memory controller: Xilinx Corporation Device 923f
81:00.3 Memory controller: Xilinx Corporation Device 933f
```

```
[xilinx@] # cat /sys/bus/pci/devices/0000:81:00.0/sriov_totalvfs
```

This command provided the maximum number of VFs, the current PF can be assigned with.

- ✓ Assign the VFs to a PF:

```
[xilinx@] # echo x > /sys/bus/pci/devices/0000:81:00.0/sriov_numvfs
```

This command allows x number of VFs to get assigned to the current PF.

Once the VFs are assigned to the PF, lspci lists the newly instantiated VF devices.

- ✓ Attaching the VFs to VM:
The newly instantiated VFs can now be attached to the VMs, if VM is installed on the Host

3.8 *Running the VF on Virtual Machines*

- Create a new VM using virt-manager or any similar tools

- Insert qdma driver in host machine

```
[xilinx@] # insmod qdma.ko
```

- Allocate the number of Qs for VF by writing into qmax_vfs on the master_pf device

```
[xilinx@] # echo 1000 > /sys/bus/pci/devices/<master_pf_device>/qdma/qmax_vfs
```

- Instantiate VFs on host side

```
[xilinx@] # echo 1 > /sys/bus/pci/devices/<master_pf_device>/sriov_numvfs
```

- Remove any qdma_vf driver if present in host side

```
[xilinx@] # rmmod qdma_vf
```

- Attach the required VF device to VM the using virt-manager Add Hardware > PCI Host Device > Xilinx Corporation device. For configuration using virsh commands, please refer section [2.9.1](#)

- Start the VM

- Once the system is booted, Insert the vf driver on VM

```
[xilinx@] # insmod qdma_vf
```

- Set the required number of Qs for the VF using vf qmax interface

```
[xilinx@] # echo 10 > /sys/bus/pci/devices/<vf_id>/qdma/qmax
```

- Now the system is ready to perform the transfers.

3.8.1 Setting up the VM using virsh commands

The virsh program is the main interface for managing virsh guest domains. The program can be used to manage the VMs in a domain, including tasks like create, pause, shutdown, list etc. It can also be used for attaching/detaching host side peripherals to the VMs.

QDMA Linux Kernel Reference Driver User Guide



Once the VM is created, attach the device using virsh attach command

Find out the bus/slot/function for the VF device to attach

```
[ root ] lspci |grep -i xilinx
05:00.0 Memory controller: Xilinx Corporation Device 903f
05:00.1 Memory controller: Xilinx Corporation Device 913f
05:00.2 Memory controller: Xilinx Corporation Device 923f
05:00.3 Memory controller: Xilinx Corporation Device 933f
05:00.4 Memory controller: Xilinx Corporation Device a03f
05:00.5 Memory controller: Xilinx Corporation Device a03f
05:00.6 Memory controller: Xilinx Corporation Device a03f
05:00.7 Memory controller: Xilinx Corporation Device a03f
```

Get the corresponding virsh nodes for the Xilinx VF devices using the bus/slot/function obtained in the lspci command

```
[ root ] virsh nodedev-list --cap pci | grep 05
pci_0000_05_00_0
pci_0000_05_00_1
pci_0000_05_00_2
pci_0000_05_00_3
pci_0000_05_00_4
pci_0000_05_00_5
pci_0000_05_00_6
pci_0000_05_00_7
```

The **nodedev-dumpxml** command list the corresponding xml for the virsh node and get the related information for the node using

```
[ root ] virsh nodedev-dumpxml pci_0000_05_00_5
<device>
  <name>pci_0000_05_00_5</name>
  <path>/sys/devices/pci0000:00/0000:00:03.0/0000:03:00.0/0000:04:10.0/0000:05:00.5</path>
  <parent>pci_0000_04_10_0</parent>
  <capability type='pci'>
    <domain>0</domain>
    <bus>5</bus>
    <slot>0</slot>
    <function>5</function>
    <product id='0xa03f' />
```

```

<vendor id='0x10ee'>Xilinx Corporation</vendor>
<capability type='phys_function'>
  <address domain='0x0000' bus='0x05' slot='0x00' function='0x0' />
</capability>
<iommuGroup number='44'>
  <address domain='0x0000' bus='0x05' slot='0x00' function='0x5' />
</iommuGroup>
<numa node='0' />
<pci-express>
  <link validity='cap' port='0' speed='8' width='16' />
  <link validity='sta' width='0' />
</pci-express>
</capability>
</device>

```

Once the details of the node are available, edit the configuration of the VM using `virsh edit` command. We can either manually edit the file or use `virsh` compatible xml files to attach/detach the device. This document is assuming the manual editing of the `virsh` XML configuration file for a VM.

In the `virsh` XML configuration file, `address domain` and `address type` fields represents the following information

Address domain - In host what is the bus/slot/function of the device which must be assigned to the VM. Use `'lspci'` output to figure out the bus/slot/function to use in the respective fields

Address type – In the target VM, what should be the bus/slot/function for the device. Make sure that it doesn't conflict with other entries in the configuration. If it conflicts, VM instantiation will fail with respective error messages.

```

[ root ] virsh edit vm2-ubuntu18.04
<hostdev mode='subsystem' type='pci' managed='yes'>
<source>
  <address domain='0x0000' bus='0x05' slot='0x00' function='0x4' />
</source>
<address type='pci' domain='0x0000' bus='0x00' slot='0x08' function='0x0' />
</hostdev>

```

For adding one more VF to VM, edit the configuration file and add the new entry, ensuring the address type bus/slot/function is properly configured

```

[ root ] virsh edit vm2-ubuntu18.04

<hostdev mode='subsystem' type='pci' managed='yes'>

```

```
<source>
  <address domain='0x0000' bus='0x05' slot='0x00' function='0x4' />
</source>
<address type='pci' domain='0x0000' bus='0x00' slot='0x08' function='0x0' />
</hostdev>
<hostdev mode='subsystem' type='pci' managed='yes'>
<source>
  <address domain='0x0000' bus='0x05' slot='0x00' function='0x5' />
</source>
<address type='pci' domain='0x0000' bus='0x00' slot='0x19' function='0x0' />
</hostdev>
```

Start the VM using virsh start <VM name>

```
[ root ] virsh start vm2-ubuntu18.04
```

Verify the running status of the VM using virsh list command

```
[ root ] virsh list --all
```

Id	Name	State

3	vm2-ubuntu18.04	running

Once started, you can login to the VM using the ssh command

```
[ root ] ssh xilinx@$vm2
```

Once logged into the VM, ensure that the device is properly attached.

```
xilinx@vm4-ubuntu:~$ lspci |grep -i xilinx
00:08.0 Memory controller: Xilinx Corporation Device a03f
00:19.0 Memory controller: Xilinx Corporation Device a03f
```



3.9 *Un-installing the QDMA Driver modules*

Standard Linux commands should be used to uninstall the.

- Uninstall the kernel module.

```
[xilinx@] # make uninstall
```


4 Appendix 1 – User Application “dmactl” command options

dmactl support device management commands and queue management commands. This section describes the details of each option provided for dmactl commands.

Format	Parameter	Range
q add idx <N>	Queue index	default for pf: 0 – 511 default for vf: 0 -7 If in case the number of queues per pf/vf are configured differently, the range will be changed such that 0 – (max range configured-1)
q add list <start_index> <num_Qs>	start_index: starting queue number in the range <num_Qs>: Ending queue number in the range	Same as above
[mode <mm st>]	mm or st	mm: memory mapped mode st: streaming mode
[dir <h2c c2h bi>]	h2c or c2h or bi	h2c: host to card c2h: card to host bi: both h2c and c2h
[idx_ringsz <0:15>]	Ring size	Ring size is an enum number which allows values from 0 -15. 16 different ring sizes can be configured for QDMA sub system
[idx_bufsz <0:15>]	Buffer Size	Buffer size is an enum number which allows values from 0 -15. 16 different buffer sizes can be configured for QDMA sub system
[idx_tmr <0:15>]	Timer index	The reference timer is based on the timer tick which is an enum number and allows 0 -15 values. The timer_idx in the WRB Context is the index to the 16 QDMA_C2H_TIMER_CNT registers. Each queue can choose its own

		timer_idx.
[idx_cntr <0:15>]	Counter index	Counter index from 0 to 15
[trigmode <every usr_cnt usr usr_tmr dis>]	Trigger mode	Trigger Mode: Disable:0 Any:1 Timer:2 Counter:3 Combo:4 User:5
[cmptsz <0 1 2 3>]	Completion entry size	Completion Descriptor Size: 8B:0 16B:1 32B:2 64B:3
[sw_desc_sz <3>]	Software descriptor size	Software Descriptor Size: 64B:3
[desc_bypass_en]	Enable the descriptor bypass mode	In cache bypass mode, a queue fetched descriptor is sent to user logic. User logic is then responsible for delivering the packet and associated descriptors in simple bypass interface. This option is applicable for Streaming C2H only.
[pfetch_bypass_en]	Enable the simple bypass mode.	Enable the simple bypass mode
[pfetch_en]	Enable Prefetch	When the prefetch is enabled, the prefetch engine will prefetch the descriptors from the descriptor fetch engine at the first time it fetches the descriptors for that queue. The number of descriptors that it can prefetch is defined in the registers
[dis_cmpl_status]	Disable completion status	This option allows the user to disable the completions
[dis_cmpl_status_acc]	Disable completion status accumulation	Completion status accumulation allows the completions to be triggered back after certain number of descriptors being processed. Default value is 4 descriptors. By disabling this, completions are triggered for every descriptor being processed

[dis_cmpl_status_pend_chk]	Disable completion status pending check	This option disables the completion status pending check.
[c2h_udd_en]	Enable user defined data	This option allows the user to enable the user defined data to be embedded in streaming C2H mode.
[dis_fetch_credit]	Disable fetch credit	The number of descriptors fetched will be qualified by the number of credits given to the queue. Set to 1 for C2H ST by default. This option allows to disable the fetch credit.
[dis_cmpl_status]	Disable completion status	Disable completion status
[cmpl_ovf_dis]	Disable completion ring overflow check	Disable completion ring overflow check
[c2h_cmpl_intr_en]	Enable ST C2H completion interrupts	This option allows to enable the completion interrupts
[desc <x> <y>]	Descriptor indexes from <x> to <y>	This option allows to dump the descriptors from index <x> to index <y>
[cmpt <x> <y>]	Completion indexes from <x> to <y>	This option allows to dump the completion descriptors from index <x> to index <y>
[dmap <Q> <N>]	dump dmap registers	dump dmap registers if dmap is specified. specify dmap range to dump: Q=queue, N=num of queues
udd idx <N>	User define data index	dump the user defined data received. This is applicable for ST C2H only
bar <N>	Bar index	QDMA IP has 3 bars 0: config bar 1: bypass bar 2: user bar
vector <N>	Vector index	QDMA IP supports 2K interrupt vectors and <N> is an index of the interrupt vector to be used.

5 Appendix 2 – dmautils tool

QDMA Linux driver provides character device interface for standalone IP testing. A char device is created by Linux driver for each queue pair of QDMA IP that is added to a function

It provides IO interface using following function pointers provided in f_ops structure of kernel char device driver.

- read
- write
- aio_read
- aio_write
- aio_read_iter
- aio_write_iter'

Standard IO tools like 'fio' can be used for performing IO operations using the char device interface. These standard tools pose a challenge of not being able to keep the driver/ HW busy enough while doing performance testing as they are limited to sending / receiving 1 packet at a time and wait for the processing of the packet to complete. This limitation cannot be overcome with standard tools because, if an application provides a buffer for read/write DMA operation, unless the driver confirms the completion the application cannot free the allocated buffer. The true potential of HW and driver can only be tested when application is able to send / receive enough data at higher throughput to keep the driver and HW busy.

To overcome the above said limitation, an asynchronous IO capable tool is required to provide the buffers for DMA operation continuously and free the buffers only when driver notifies the application of the completion corresponding to the IO submitted.

This can also be achieved with fio tool, but, if we want to keep the driver and HW busy, application needs to continuously submit IO requests while polling for the completion parallelly, which is not done in fio.

This can be achieved by leveraging the asynchronous functionality provided by libaio library. Using libaio, application can submit IO request to the driver and driver returns the call immediately and notify of completion separately. Application can then poll for the completion and free the buffer upon receiving the completion.

'dmautils' tool developed by Xilinx, specifically for QDMA, tries to accomplish this by continuously submitting IO requests while another thread continuously polls for completion events of IOs submitted. This tool is capable of following features:

- Highly configurable

It enables the user to configure the following

- Number of PFs and queues to do IO testing on
- Number of threads that poll on each char device
- Duration for which continuous IOs need to be done on each char device
- Dump queue context and registers on completion of IO
- Set number of packets to be sent per IO
- Set packet size
- Do q add/start/stop/delete on all queues on which IOs are required to be done
- Set queue mode (ST/MM) and direction (H2C/C2H) of each queue for IO operations

- Set some of queue configurations (pre-fetch/ring size/ST C2H completion size, timer threshold, counter threshold and trigger mode) while doing 'q start'.
- Set zero copy for ST C2H performance testing
- Performs read and write unidirectional calls and bidirectional simultaneous calls
- Supports the interface to submit IO requests from multiple threads on a single char device interface
- Calculates the number of packets for which completion is received in the time duration specified though configuration file. This gives us the number of packets per second and in turn the throughput.

Note: dmautils tool uses zero buffers and does not do any data validation as this tool is currently targeted mainly for performance testing only.

Usage of dmautils tool

```
[xilinx@] # dmautils -c "config_file"
```

Sample config files are available in "linux-kernel/tools/config/dmautils_config".

Below table covers all the available configuration parameters that can be provided to dmautils tool is given below:

Config Parameter	Example Value	Description
mode	st	Queue mode
dir	c2h	Queue direction
pf_range	0:0	Range of PFs to be used
q_range	0:7	Range of queues to be used
wb_acc	5	Writeback accumulation value. The writeback accumulation will happen for $2^{(\text{value} + 1)}$
dump_en	0	Enable logging of queue context, register dump and lspci output for every queue for each IO size performed
tmr_idx	5	Timer index to be selected for ST C2H writeback accumulation
cntr_idx	6	Counter index to be selected for ST C2H writeback accumulation
trig_mode	cntr_tmr	Trigger mode for ST C2H writeback update
pfetch_en	1	Flag to enable prefetch on the queues
wrbsz	1	Completion entry size
rngidx	5	Ring size index selection
runtime	30	Time duration for which the IO should be performed for each IO size
num_threads	4	Number of threads that should do IOs on each queue simultaneously
num_pkt	64	Number of packets to be IO'ed at once
pkt_sz	64	Packet size
pci_bus	17	PCI bus number
pci_device	00	PCI device number

Table 5-1: dmautils tool configuration options

6 Appendix 3 – Release Directory Structure

6.1 SW Directory

The entire software source is under https://github.com/Xilinx/dma_ip_drivers subdirectory **QDMA/linux-kernel/** folder

Directory	Description
linux-kernel	Top-level directory for QDMA Linux SW driver, example application, documents and tools software
docs/	Documentation for the QDMA Linux Driver
drv/	Provides the interfaces to manage the underlined PCIe device and provide character interface to control the QDMA IP
libqdma/	QDMA library, used by the source in drv/
user/	User space application to configure and control the QDMA IP
tools/	Tools to perform DMA operations
Makefile	Make file to compile the Linux QDMA Driver

Table 6-1: SW Directory

7 Appendix 4 – Doxygen tool usage for document generation

For generating the pdf documentation from the source code using Doxygen tool, the following software is required to install on the host system

- Doxygen
- texlive-latex-base

Doxygen configuration file(Doxyfile) is provided in the release package which has the necessary settings for document generation.

7.1 *Steps for document generation*

- Open Linux terminal & change directory to QDMA/linux-kernel/docs/

```
[xilinx@] # cd QDMA/linux-kernel/docs
```

- run the below command

```
[xilinx@docs] # doxygen Doxyfile
```

- Tool generates the documentation in latex format at qdma/latex. Change to the directory

```
[xilinx@docs] # cd qdma/latex
```

- Execute make command to build and generate pdf document

```
[xilinx@latex] # make
```

- PDF document gets generated in the same directory with name “refman.pdf”